our chefs!

Eleni Skylaki

María García Guerrero

Mark Swaringen

Matthaeus Hilpold

BE CAREFUL!
EXTRA SPICY
HOMEMADE AUGMENTATION

SECRET TRAINING RECIPE

FIRST BATCH

Start with a modern CNN architecture, adding pre-trained weights.
Next, freeze the model and rebuild the top with Global Avg Pooling, Batch Normalization, Dropout, and finally a dense layer with Softmax activation, followed by a training session.

SECOND BATCH

Now unfreeze the entire model (with the exception of Batch Normalization) and train again. Voilà, dinner is served!

100% Fresh MODELS

FOR KERAS' LOVERS

NOISY-STUDENT WEIGHTS
CHEF'S CHOICE TRANSFER-LEARNING

RESNET & EFFICIENTNETB0-B5
PLUS EXTRA 20 LAYERS RETRAINED!

CUSTOM LR-SCHEDULER
RAMPUP WITH STEP-DECAY/COSINE

ENSEMBLE BOOST!

TASTY PREDICTIONS!

Not Hot Dog v2

DIVINE MACHINE LEARNING

ESTABLISHED IN 2020

our menus!

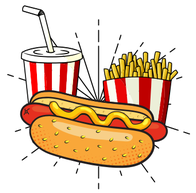

EFFICIENTNET B5 — 68%

50 epochs (with early stopping), 456x456 img resolution, batch size 32, cosine learning decay, 0.5 top-dropout rate, Noisy-Student weights, Adam optimizer, categorical crossentropy activation, softmax loss, avg pooling & batch normalization layers.

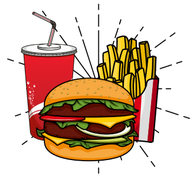

EFFICIENTNET B0 — 66%

Kid menu version of B5: smaller net, 224x224 img resolution, batch-size 18 and step-decay learning.

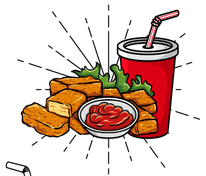

RESNET 152 — 58%

Traditional dish: 224x224 img resolution, weights from imagenet, 11 epochs, batch-size 32, relu activation for the dense layer, 0.2 dropout, softmax classifier, Adam optimizer.

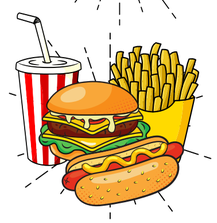

"TOUS ENSEMBLE" FAMILY MENU — 70%

Max voting system: the most voted label among all models makes the final prediction. If the models disagree, the best model's prediction is chosen.

GIVE YOUR MENU A TWIST!

ALBUMENTATIONS
RGB shift, blur, saturation and much more....

MORE GENERALIZATION
Random flip, crop, scale...